

# CS145 Homework 1

**Important Note:** HW1 is due on **11:59 PM PT, Oct 19 (Monday, Week 3)**. Please submit through GradeScope (you will receive an invite to Gradescope for CS145 Fall 2020.).

## Print Out Your Name and UID

Name: Terry Ye, UID: 004757414

## Before You Start

You need to first create HW1 conda environment by the given `cs145hw1.yml` file, which provides the name and necessary packages for this tasks. If you have `conda` properly installed, you may create, activate or deactivate by the following commands:

```
conda env create -f cs145hw1.yml
conda activate hw1
conda deactivate
```

OR

```
conda env create --name NAMEOFOURCHOICE -f cs145hw1.yml
conda activate NAMEOFOURCHOICE
conda deactivate
```

To view the list of your environments, use the following command:

```
conda env list
```

More useful information about managing environments can be found [here](https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html) (<https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>).

You may also quickly review the usage of basic Python and Numpy package, if needed in coding for matrix operations.

In this notebook, you must not delete any code cells in this notebook. If you change any code outside the blocks that you are allowed to edit (between `START/END YOUR CODE HERE`), you need to highlight these changes. You may add some additional cells to help explain your results and observations.

```
In [1]: import numpy as np
import pandas as pd
import sys
import random as rd
import matplotlib.pyplot as plt
%load_ext autoreload
%autoreload 2
```

If you can successfully run the code above, there will be no problem for environment setting.

# 1. Linear regression

This workbook will walk you through a linear regression example.

```
In [2]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
# As a sanity check, we print out the size of the training data (1000, 100) and test data (100, 100)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)
```

Training data shape: (1000, 100)  
Training labels shape: (1000,)

```
In [3]: def train_and_predict(lm, beta):
    train_predict_y = lm.predict(lm.train_x, beta)
    test_predict_y = lm.predict(lm.test_x, beta)
    training_error = lm.compute_mse(train_predict_y, lm.train_y)
    testing_error = lm.compute_mse(test_predict_y, lm.test_y)
    return training_error, testing_error
```

## 1.1 Closed form solution

In this section, complete the `getBeta` function in `linear_regression.py` which use the closed form solution of  $\hat{\beta}$ .

Train your model by using `lm.train('0')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [4]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# START YOUR CODE HERE #
#####
beta = lm.train('0')
training_error, testing_error = train_and_predict(lm, beta)
#####
# END YOUR CODE HERE #
#####
print('Training error is: ', training_error)
print('Testing error is: ', testing_error)
```

Learning Algorithm Type: 0  
Training error is: 0.08693886675396784  
Testing error is: 0.11017540281675803

## 1.2 Batch gradient descent

In this section, complete the `getBetaBatchGradient` function in `linear_regression.py` which compute the gradient of the objective fuction.

Train you model by using `lm.train('1')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [5]: """add import here so can just run this cell alone when update the source code"""
from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####
beta = lm.train('1')
training_error, testing_error = train_and_predict(lm, beta)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_error)
print('Testing accuracy is: ', testing_error)
```

```
Learning Algorithm Type: 1
Training accuracy is: 0.08693895533150824
Testing accuracy is: 0.11016592170824556
```

## 1.3 Stochastic gadient descent

In this section, complete the `getBetaStochasticGradient` function in `linear_regression.py` , which use an estimated gradient of the objective function.

Train you model by using `lm.train('2')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [6]: """add import here so can just run this cell alone when update the source code"""
from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####
beta = lm.train('2')
training_error, testing_error = train_and_predict(lm, beta)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_error)
print('Testing accuracy is: ', testing_error)
```

Learning Algorithm Type: 2  
 Training accuracy is: 0.09337211730536127  
 Testing accuracy is: 0.11828923255919811

```
In [7]: """Normalize data"""
from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
lm.normalize()
for train_type in ['0', '1', '2']:
    beta = lm.train(train_type)
    training_error, testing_error = train_and_predict(lm, beta)
    print('Training error for ' + train_type + ' = ', training_error )
    print('Testing error for ' + train_type + ' = ', testing_error )
```

Learning Algorithm Type: 0  
 Training error for 0 = 0.08693886675396784  
 Testing error for 0 = 0.11017540281675804  
 Learning Algorithm Type: 1  
 Training error for 1 = 0.10024222412685467  
 Testing error for 1 = 0.14039955364152018  
 Learning Algorithm Type: 2  
 Training error for 2 = 0.0869952008695027  
 Testing error for 2 = 0.1099647262725722

## Questions:

1. Compare the MSE on the testing dataset for each version. Are they the same? Why or why not?
2. Apply z-score normalization for eachh featrure and comment whether or not it affect the three algorithm.
3. Ridge regression is adding an L2 regularization term to the original objective function of mean squared error. The objective function become following:

$$J(\beta) = \frac{1}{2n} \sum_i (x_i^T \beta - y_i)^2 + \frac{\lambda}{2n} \sum_j \beta_j^2,$$

where  $\lambda \geq 0$ , which is a hyper parameter that controls the trade off. Take the derivative of this provided objective function and derive the closed form solution for  $\beta$ .

## Your answer here:

1. The testing errors are not the same and slightly different. The Closed Form and Batch Gradient performed similarly while the Stochastic Gradient performed a bit worse. I think Closed Form and Batch Gradient all reached the optimum and the little difference is because Batch Gradient has learning rate which prevents it from achieving the absolute best in limited iterations. The Stochastic Gradient performed worse maybe because the learning rate is smaller and it converges much slower.

2. It does not affect the Closed Form algorithm but affects Batch Gradient and Stochastic Gradient algorithms.

3. The matrix form of objective function can be written as  $\frac{1}{2n} (X\beta - y)^T (X\beta - y) + \frac{\lambda}{2n} \beta^T \beta$

Taking derivative of the matrix form as 0:  $J(\beta)' = (X^T X \beta - X^T y)/n + \frac{\lambda}{n} \beta = 0$

So  $(X^T X + \lambda I)\beta = X^T y$ , the closed form solution of  $\beta$  is

$$\beta = (X^T X + \lambda I)^{-1} X^T y$$

## 2. Logistic regression

This workbook will walk you through a logistic regression example.

```
In [8]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
# As a sanity check, we print out the size of the training data (1000, 5) and test data (1000, 5)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape: ', lm.train_y.shape)
```

Training data shape: (1000, 5)

Training labels shape: (1000,)

```
In [9]: def train_and_predict_logistic(lm, beta):
    train_predict_y = lm.predict(lm.train_x, beta)
    test_predict_y = lm.predict(lm.test_x, beta)
    training_accuracy = lm.compute_accuracy(train_predict_y, lm.train_y)
    testing_accuracy = lm.compute_accuracy(test_predict_y, lm.test_y)
    return training_accuracy, testing_accuracy
```

### 2.1 Batch gradient descent

In this section, complete the `getBeta_BatchGradient` in `logistic_regression.py`, which compute the gradient of the log likelihood function.

Complete the `compute_avglogL` function in `logistic_regression.py` for sanity check.

Train your model by using `lm.train('0')` function.

And print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```
In [10]: """add import here so can just run this cell alone when update the source code"""
from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####
lm.normalize()
beta = lm.train('0')
training_accuracy, testing_accuracy = train_and_predict_logistic(lm, beta)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)
```

```
average logL for iteration 0: -0.48427563333314566
average logL for iteration 1000: -0.46010037535085324
average logL for iteration 2000: -0.46010037535085324
average logL for iteration 3000: -0.46010037535085324
average logL for iteration 4000: -0.46010037535085324
average logL for iteration 5000: -0.46010037535085324
average logL for iteration 6000: -0.46010037535085324
average logL for iteration 7000: -0.46010037535085324
average logL for iteration 8000: -0.46010037535085324
average logL for iteration 9000: -0.46010037535085324
Training avgLogL: -0.46010037535085324
Training accuracy is: 0.797
Testing accuracy is: 0.7534791252485089
```

## 2.2 Newton Raphhson

In this section, complete the `getBeta_Newton` in `logistic_regression.py`, which make use of both first and second derivative.

Train your model by using `lm.train('1')` function.

Print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```
In [11]: """add import here so can just run this cell alone when update the source code"""
from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-t
training_accuracy= 0
testing_accuracy= 0
#=====#
# START YOUR CODE HERE #
#=====#
lm.normalize()
beta = lm.train('1')
training_accuracy, testing_accuracy = train_and_predict_logistic(lm, beta)
#=====#
# END YOUR CODE HERE #
#=====#
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)
```

```
average logL for iteration 0: -0.5564001859888966
average logL for iteration 500: -0.46010037535085324
average logL for iteration 1000: -0.46010037535085324
average logL for iteration 1500: -0.46010037535085324
average logL for iteration 2000: -0.46010037535085324
average logL for iteration 2500: -0.46010037535085324
average logL for iteration 3000: -0.46010037535085324
average logL for iteration 3500: -0.46010037535085324
average logL for iteration 4000: -0.46010037535085324
average logL for iteration 4500: -0.46010037535085324
average logL for iteration 5000: -0.46010037535085324
average logL for iteration 5500: -0.46010037535085324
average logL for iteration 6000: -0.46010037535085324
average logL for iteration 6500: -0.46010037535085324
average logL for iteration 7000: -0.46010037535085324
average logL for iteration 7500: -0.46010037535085324
average logL for iteration 8000: -0.46010037535085324
average logL for iteration 8500: -0.46010037535085324
average logL for iteration 9000: -0.46010037535085324
average logL for iteration 9500: -0.46010037535085324
Training avgLogL: -0.46010037535085324
Training accuracy is: 0.797
Testing accuracy is: 0.7534791252485089
```

## Questions:

1. Compare the accuracy on the testing dataset for each version. Are they the same? Why or why not?
2. Regularization. Similar to linear regression, an regularization term could be added to logistic regression. The objective function becomes following:

$$J(\beta) = -\frac{1}{n} \sum_i (y_i x_i^T \beta - \log(1 + \exp\{x_i^T \beta\})) + \lambda \sum_j \beta_j^2,$$

where  $\lambda \geq 0$ , which is a hyper parameter that controls the trade off. Take the derivative  $\frac{\partial J(\beta)}{\partial \beta_j}$  of this provided objective function and provide the batch gradient descent update.

### Your answer here:

1. They are the same because both methods achieve the optimum for log likelihood and the result beta is unique.
2. The derivative of new objective function is the derivative of original J + derivate of regularization term.

$$\frac{\partial J(\beta)}{\partial \beta_j} = \sum_{i=1}^N x_{ij}(y_i - p_i(\beta)) + 2\lambda\beta_j$$

The batch gradient update is

$$\beta^{new} = \beta^{old} + \eta v_j = \sum_{i=1}^N x_{ij}(y_i - p_i(\beta)) + 2\lambda\beta_j$$

## 2.3 Visualize the decision boundary on a toy dataset

In this subsection, you will use the same implementation for another small dataset with each datapoint  $x$  with only two features ( $x_1, x_2$ ) to visualize the decision boundary of logistic regression model.

```
In [12]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression(verbose = False)
lm.load_data('./data/logistic-regression-toy.csv', './data/logistic-regression-toy')
# As a sanity check, we print out the size of the training data (99,2) and training labels
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)
```

```
Training data shape: (99, 2)
Training labels shape: (99,)
```

In the following block, you can apply the same implementation of logistic regression model (either in 2.1 or 2.2) to the toy dataset. Print out the  $\hat{\beta}$  after training and accuracy on the train set.



```
In [13]: training_accuracy= 0
#=====#
#  STRART YOUR CODE HERE  #
#=====#
beta = lm.train('0')
train_predict_y = lm.predict(lm.train_x, beta)
training_accuracy = lm.compute_accuracy(train_predict_y, lm.train_y)
print(beta)
#=====#
#  END YOUR CODE HERE  #
#=====#
print('Training accuracy is: ', training_accuracy)
```

```
Training avgLogL:  -0.329147431295712
[-2.6205116   0.76037154  1.17194674]
Training accuracy is:  0.8888888888888888
```

Next, we try to plot the decision boundary of your learned logistic regression classifier. Generally, a decision boundary is the region of a space in which the output label of a classifier is ambiguous. That is, in the given toy data, given a datapoint  $x = (x_1, x_2)$  on the decision boundary, the logistic regression classifier cannot decide whether  $y = 0$  or  $y = 1$ .

## Question

Is the decision boundary for logistic regression linear? Why or why not?

## Your answer here:

Yes it is linear as there are only two variables involved with only first indexes.

Draw the decision boundary in the following cell. Note that the code to plot the raw data points are given. You may need `plt.plot` function (see [here](https://matplotlib.org/tutorials/introductory/pyplot.html) (<https://matplotlib.org/tutorials/introductory/pyplot.html>)).

```

In [14]: # scatter plot the raw data
df = pd.concat([lm.train_x, lm.train_y], axis=1)
groups = df.groupby("y")
for name, group in groups:
    plt.plot(group["x1"], group["x2"], marker="o", linestyle="", label=name)


# plot the decision boundary on top of the scattered points
#=====#
# STRART YOUR CODE HERE #
#=====#

x1_values = [i for i in np.arange(0,9,0.01)]
x2_values = []

# boundary: x1*beta[1] + x2*beta[2] + beta[0] = 0
for x1 in x1_values:
    x2 = (-beta[0] - x1*beta[1]) / beta[2]
    x2_values.append(x2)

plt.plot(x1_values, x2_values)
#=====#
# END YOUR CODE HERE #
#=====#
plt.show()

```



## End of Homework 1 :)

After you've finished the homework, please print out the entire `ipynb` notebook and two `py` files into one PDF file. Make sure you include the output of code cells and answers for questions. Prepare submit it to GradeScope.

```

1 import pandas as pd
2 import numpy as np
3 import sys
4 import random as rd
5
6 #insert an all-one column as the first column
7 def addAllOneColumn(matrix):
8     n = matrix.shape[0] #total of data points
9     p = matrix.shape[1] #total number of attributes
10
11     newMatrix = np.zeros((n,p+1))
12     newMatrix[:,1:] = matrix
13     newMatrix[:,0] = np.ones(n)
14
15     return newMatrix
16
17 # Reads the data from CSV files, converts it into Dataframe and returns x and y
18 # dataframes
19 def getDataframe(filePath):
20     dataframe = pd.read_csv(filePath)
21     y = dataframe['y']
22     x = dataframe.drop('y', axis=1)
23     return x, y
24
25 # train_x and train_y are numpy arrays
26 # function returns value of beta calculated using (0) the formula  $\beta = (X^T X)^{-1} (X^T Y)$ 
27 def getBeta(train_x, train_y):
28     n = train_x.shape[0] #total of data points
29     p = train_x.shape[1] #total number of attributes
30
31     beta = np.zeros(p)
32     #=====#
33     # STRART YOUR CODE HERE #
34     #=====#
35
36     x_transpose = np.transpose(train_x)
37     left_operand = np.linalg.inv(np.matmul(x_transpose, train_x))
38     right_operand = np.matmul(x_transpose, train_y)
39     beta = np.matmul(left_operand, right_operand)
40
41     #=====#
42     # END YOUR CODE HERE #
43     #=====#
44     return beta
45
46 # train_x and train_y are numpy arrays
47 # lr (learning rate) is a scalar
48 # function returns value of beta calculated using (1) batch gradient descent
49 def getBetaBatchGradient(train_x, train_y, lr, num_iter):
50     beta = np.random.rand(train_x.shape[1])
51
52     n = train_x.shape[0] #total of data points
53     p = train_x.shape[1] #total number of attributes
54
55     beta = np.random.rand(p)
56     #update beta iteratively
57     for iter in range(0, num_iter):
58         deriv = np.zeros(p)

```

```

59     for i in range(n):
60         #=====#
61         # STRART YOUR CODE HERE #
62         #=====#
63
64         # gradient of OLS(1/2 (yi-xi)^2)
65         cur_deriv = train_x[i] * (np.dot(beta, np.transpose(train_x[i])) -
train_y[i])
66         deriv = np.add(deriv, cur_deriv)
67
68         #=====#
69         #   END YOUR CODE HERE   #
70         #=====#
71     deriv = deriv / n
72     beta = beta - deriv.dot(lr)
73     return beta
74
75 # train_x and train_y are numpy arrays
76 # lr (learning rate) is a scalar
77 # function returns value of beta calculated using (2) stochastic gradient descent
78 def getBetaStochasticGradient(train_x, train_y, lr):
79     n = train_x.shape[0] #total of data points
80     p = train_x.shape[1] #total number of attributes
81
82     beta = np.random.rand(p)
83
84     epoch = 100
85     for iter in range(epoch):
86         indices = list(range(n))
87         rd.shuffle(indices)
88         for i in range(n):
89             idx = indices[i]
90             #=====#
91             # STRART YOUR CODE HERE #
92             #=====#
93
94             # use np.multiply instead of * here to avoid overflow
95             # needs to multiply lr first to avoid overflow
96             coefficient = lr * (train_y[idx] - np.dot(np.transpose(train_x[idx]),
beta))
97             cur_update = np.multiply(coefficient, train_x[idx])
98             beta = np.add(beta, cur_update)
99
100            #=====#
101            #   END YOUR CODE HERE   #
102            #=====#
103    return beta
104
105
106 # Linear Regression implementation
107 class LinearRegression(object):
108     # Initializes by reading data, setting hyper-parameters, and forming linear model
109     # Forms a linear model (learns the parameter) according to type of beta (0 -
closed form, 1 - batch gradient, 2 - stochastic gradient)
110     # Performs z-score normalization if z_score is 1
111     def __init__(self, lr=0.005, num_iter=1000):
112         self.lr = lr
113         self.num_iter = num_iter
114         self.train_x = pd.DataFrame()
115         self.train_y = pd.DataFrame()

```

```
116     self.test_x = pd.DataFrame()
117     self.test_y = pd.DataFrame()
118     self.algType = 0
119     self.isNormalized = 0
120
121     def load_data(self, train_file, test_file):
122         self.train_x, self.train_y = getDataframe(train_file)
123         self.test_x, self.test_y = getDataframe(test_file)
124
125     def normalize(self):
126         # Applies z-score normalization to the dataframe and returns a normalized
dataframe
127         self.isNormalized = 1
128         means = self.train_x.mean(0)
129         std = self.train_x.std(0)
130         self.train_x = (self.train_x - means).div(std)
131         self.test_x = (self.test_x - means).div(std)
132
133         # Gets the beta according to input
134     def train(self, algType):
135         self.algType = algType
136         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one column
as the first column
137         print('Learning Algorithm Type: ', algType)
138         if(algType == '0'):
139             beta = getBeta(newTrain_x, self.train_y.values)
140             #print('Beta: ', beta)
141
142         elif(algType == '1'):
143             beta = getBetaBatchGradient(newTrain_x, self.train_y.values, self.lr,
self.num_iter)
144             #print('Beta: ', beta)
145         elif(algType == '2'):
146             # change learning rate to 0.0005 to converge
147             beta = getBetaStochasticGradient(newTrain_x, self.train_y.values, 0.0005)
148             #print('Beta: ', beta)
149         else:
150             print('Incorrect beta_type! Usage: 0 - closed form solution, 1 - batch
gradient descent, 2 - stochastic gradient descent')
151
152
153         return beta
154
155     # Predicts the y values on given data and learned beta
156     def predict(self, x, beta):
157         newTest_x = addAllOneColumn(x)
158         self.predicted_y = newTest_x.dot(beta)
159         return self.predicted_y
160
161
162     # predicted_y and y are the predicted and actual y values respectively as numpy
arrays
163     # function returns the mean squared error (MSE) value for the test dataset
164     def compute_mse(self, predicted_y, y):
165         mse = np.sum((predicted_y - y)**2)/predicted_y.shape[0]
166         return mse
167
168
169
```

```

1  # -*- coding: utf-8 -*-
2
3  import pandas as pd
4  import numpy as np
5  import sys
6  import random as rd
7
8  #insert an all-one column as the first column
9  def addAllOneColumn(matrix):
10     n = matrix.shape[0] #total of data points
11     p = matrix.shape[1] #total number of attributes
12
13     newMatrix = np.zeros((n,p+1))
14     newMatrix[:,0] = np.ones(n)
15     newMatrix[:,1:] = matrix
16
17
18     return newMatrix
19
20 # Reads the data from CSV files, converts it into Dataframe and returns x and y
dataframes
21 def getDataframe(filePath):
22     dataframe = pd.read_csv(filePath)
23     y = dataframe['y']
24     x = dataframe.drop('y', axis=1)
25     return x, y
26
27 # sigmoid function
28 def sigmoid(z):
29     return 1 / (1 + np.exp(-z))
30
31 # compute average logL
32 def compute_avglogL(X,y,beta):
33     eps = 1e-50
34     n = y.shape[0]
35     avglogL = 0
36     #=====#
37     # STRART YOUR CODE HERE #
38     #=====#
39     total_logL = 0
40     for i in range(n):
41         # logL = yiXi^TB - log(1 + exp{xi^TB})
42         XiT_beta = np.dot(np.transpose(X[i]), beta)
43         left_oper = np.multiply(y[i], XiT_beta)
44         right_oper = np.log(1 + np.exp(XiT_beta))
45         total_logL += (left_oper - right_oper)
46     avglogL = total_logL / n
47     #=====#
48     # END YOUR CODE HERE #
49     #=====#
50     return avglogL
51
52
53 # train_x and train_y are numpy arrays
54 # lr (learning rate) is a scalar
55 # function returns value of beta calculated using (0) batch gradient descent
56 def getBeta_BatchGradient(train_x, train_y, lr, num_iter, verbose):
57     beta = np.random.rand(train_x.shape[1])
58
59     n = train_x.shape[0] #total of data points

```

```

60 p = train_x.shape[1] #total number of attributes
61
62
63 beta = np.random.rand(p)
64 #update beta iteratively
65 for iter in range(0, num_iter):
66     #=====#
67     # STRART YOUR CODE HERE #
68     #=====#
69
70     """
71     logL = yiXi^TB - log(1 + exp{xi^TB})
72     deriv = sum yixij - sum x_ij exp{b^Txo}/ (1+exp{bTxi})
73     """
74
75     deriv = np.zeros(p)
76     for j in range(p):
77         jth_deriv = 0
78         for i in range(n):
79             betaT_xi = np.dot(beta, train_x[i])
80             yi_xij = train_y[i] * train_x[i][j]
81             pi_xij = train_x[i][j] * np.exp(betaT_xi) / (1 + np.exp(betaT_xi))
82             jth_deriv = jth_deriv + yi_xij - pi_xij
83         deriv[j] += jth_deriv
84     beta = beta + np.multiply(lr, deriv)
85     #=====#
86     # END YOUR CODE HERE #
87     #=====#
88     if(verbose == True and iter % 1000 == 0):
89         avgLogL = compute_avglogL(train_x, train_y, beta)
90         print(f'average logL for iteration {iter}: {avgLogL} \t')
91     return beta
92
93 # train_x and train_y are numpy arrays
94 # function returns value of beta calculated using (1) Newton-Raphson method
95 def getBeta_Newton(train_x, train_y, num_iter, verbose):
96     n = train_x.shape[0] #total of data points
97     p = train_x.shape[1] #total number of attributes
98
99     beta = np.random.rand(p)
100     for iter in range(0, num_iter):
101         #=====#
102         # STRART YOUR CODE HERE #
103         #=====#
104
105         # calculate hessian matrix
106         # -sum X_ij X_in p_i(beta)(1 - p_i(beta))
107         hessian = np.zeros((p, p))
108         for row in range(p):
109             for col in range(p):
110                 for i in range(n):
111                     betaT_xi = np.dot(beta, train_x[i])
112                     pi_beta = np.exp(betaT_xi) / (1 + np.exp(betaT_xi))
113                     hessian[row][col] -= (train_x[i][row] * train_x[i][col] * pi_beta
114 * (1 - pi_beta))
115
116         # calculate first derivative same as getBeta_BatchGradient
117         deriv = np.zeros(p)
118         for j in range(p):
119             jth_deriv = 0

```

```

119         for i in range(n):
120             betaT_xi = np.dot(beta, train_x[i])
121             yi_xij = train_y[i] * train_x[i][j]
122             pi_xij = train_x[i][j] * np.exp(betaT_xi) / (1 + np.exp(betaT_xi))
123             jth_deriv = jth_deriv + yi_xij - pi_xij
124         deriv[j] += jth_deriv
125
126     beta = beta - np.matmul(np.linalg.inv(hessian), deriv)
127     #=====#
128     #   END YOUR CODE HERE   #
129     #=====#
130     if(verbose == True and iter % 500 == 0):
131         avgLogL = compute_avglogL(train_x, train_y, beta)
132         print(f'average logL for iteration {iter}: {avgLogL} \t')
133     return beta
134
135
136
137 # Logistic Regression implementation
138 class LogisticRegression(object):
139     # Initializes by reading data, setting hyper-parameters
140     # Learns the parameter using (0) Batch gradient (1) Newton-Raphson
141     # Performs z-score normalization if isNormalized is 1
142     # Print intermidate training loss if verbose = True
143     def __init__(self, lr=0.005, num_iter=10000, verbose = True):
144         self.lr = lr
145         self.num_iter = num_iter
146         self.verbose = verbose
147         self.train_x = pd.DataFrame()
148         self.train_y = pd.DataFrame()
149         self.test_x = pd.DataFrame()
150         self.test_y = pd.DataFrame()
151         self.algType = 0
152         self.isNormalized = 0
153
154
155     def load_data(self, train_file, test_file):
156         self.train_x, self.train_y = getDataframe(train_file)
157         self.test_x, self.test_y = getDataframe(test_file)
158
159     def normalize(self):
160         # Applies z-score normalization to the dataframe and returns a normalized
dataframe
161         self.isNormalized = 1
162         data = np.append(self.train_x, self.test_x, axis = 0)
163         means = data.mean(0)
164         std = data.std(0)
165         self.train_x = (self.train_x - means).div(std)
166         self.test_x = (self.test_x - means).div(std)
167
168     # Gets the beta according to input
169     def train(self, algType):
170         self.algType = algType
171         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one column
as the first column
172         if(algType == '0'):
173             beta = getBeta_BatchGradient(newTrain_x, self.train_y.values, self.lr,
self.num_iter, self.verbose)
174             #print('Beta: ', beta)
175

```



```
176         elif(algType == '1'):
177             beta = getBeta_Newton(newTrain_x, self.train_y.values, self.num_iter,
self.verbose)
178             #print('Beta: ', beta)
179         else:
180             print('Incorrect beta_type! Usage: 0 - batch gradient descent, 1 -
Newton-Raphson method')
181
182             train_avglogL = compute_avglogL(newTrain_x, self.train_y.values, beta)
183             print('Training avgLogL: ', train_avglogL)
184
185             return beta
186
187         # Predict on given data x with learned parameter beta
188         def predict(self, x, beta):
189             newTest_x = addAllOneColumn(x)
190             self.predicted_y = (sigmoid(newTest_x.dot(beta))>=0.5)
191             return self.predicted_y
192
193         # predicted_y and y are the predicted and actual y values respectively as numpy
arrays
194         # function returns the accuracy
195         def compute_accuracy(self,predicted_y, y):
196             acc = np.sum(predicted_y == y)/predicted_y.shape[0]
197             return acc
198
199
```